

Debugging

- finding and eliminating errors from programs

Process

1. Reproduce the bug
2. Simplify program input
3. Use a debugger to track down the origin of the problem
4. Fix the problem

Debugger

- a program that is used to run and debug other (target) programs
- advantages:
 - programmer can:
 - step through source code line by line
 - interact with and inspect program at run-time
 - if program crashes, debugger outputs where and why it crashed

GDB - GNU Debugger

- debugger for several languages
 - C, C++, Java, Objective-C, etc.
- allows you to inspect what the program is doing at a certain point during execution
- logical errors and segmentation faults are easier to find with the help of gdb

Using GDB

1. Compile program

- **normally:** `$ gcc [flags] <source files> -o <output file>`
- **debugging:** `$ gcc [other flags] -g <source files> -o <output file>`
 - `-g` option enables built-in debugging support

2. Specify program to debug

```
1 $ gdb <executable>
2 # or
```

```
3 $ gdb
4 (gdb) file <executable>
```

3. Run program

- (gdb) run [optional-arguments]

4. In GDB interactive shell

- Tab to Autocomplete
- up-arrow arrows to recall history
- help [command] to get more info about a command

5. Exit GDB

- (gdb) quit

Run-Time Errors

Segmentation fault

```
1 Program received signal SIGSEGV, Segmentation fault. 0x000000000400524 in function
(arr=0x7fffc902a270, r1=2, c1=5, r2=4, c2=6) at file.c:12
```

- line number where it crashed and parameters to the function that caused the error

Logic Error

- program will run and exit successfully

Setting Breakpoints

- used to stop running program at specific point

```
1 (gdb) break file1.c:6
2 # Program will pause when it reaches line 6 of file1.c
3
4 (gdb) break my_function
5 # Program will pause at top of my_function every time it is called
6
7 (gdb) break [position] if expression
8 # Program will pause at position only when expression evaluates to true
```

- you can view a list of all breakpoints via
 - (gdb) info|i breakpoints|break|br|b

Deleting, Disabling, and Ignoring Breakpoints

```
1 (gdb) delete [bp_number | range]
2 # Deletes specified breakpoint or range of breakpoints
3
4 (gdb) disable [bp_number | range]
5 # Temporarily disables a breakpoint or range of breakpoints
6
7 (gdb) enable [bp_number | range]
8 # Restores disabled breakpoints
9
10 # If no arguments are provided to above commands, all breakpoints are affected
11
12 (gdb) ignore bp_number iterations
13 # Instructs GDB to pass over a breakpoint without stopping a certain number of times
```

Displaying Data

```
1 (gdb) print [/format] expression
2 # Prints value of specified expression in the specified format
```

Formats

- d: decimal notation (default format for integers)
- x: hexadecimal
- o: octal
- t: binary

Resuming Execution After a Break

- when a program stops at a breakpoint, there are 4 possible kinds of gdb operations
 - **(c)ontinue**: debugger will continue executing until next breakpoint
 - **(s)tep**: debugger will continue to next source line
 - **(n)ext**: debugger will continue to next source line in the current stack frame (function)
 - **(f)inish**: debugger will resume execution until current function returns
 - execution stops immediately after program flow returns to function's caller
 - function's return value and line containing next statement are displayed


Watchpoints

- watch/observe changes to variables

```
1 (gdb) watch my_var
2 # sets a watchpoint on my_var
3 # debugger will stop the program when the value of my_var changes
4 # old and new values will be printed
5
6 (gdb) rwatch expression
7 # debugger stops the program whenever the program reads the value of any object involved in
  the evaluation of "expression"
```

Process Memory Layout

(Higher Address)



(Lower Address)

- **TEXT segment**
 - contains machine instruction to be executed
- **global variables**
 - initialized
 - uninitialized
- **heap segment**
 - dynamic memory allocation
 - malloc, free
- **stack segment**
 - push frame: function invoked
 - pop frame: function returned
 - stores
 - local variables
 - return address, registers, etc.
- **command line arguments and environment variables**

Stack Info

- a program is made up of one or more functions that interact by calling each other
- every time a function is called, an area of memory is set aside for it
 - area is called a **stack frame** and holds the following crucial info
 - storage space for all local variables
 - memory address to return to when the called function returns/ends
 - arguments, or parameters of the called function
- each function call gets its own stack frame
 - all stack frames make up the **call stack**

```
1  #include <stdio.h>
2  void first_function(void);
3  void second_function(int);
4
5  int main(void)
6  {
7      printf("hello world\n");
8      first_function();
9      printf("goodbye goodbye\n");
10
11     return 0;
12 }
13
14
15 void first_function(void)
16 {
17     int imidate = 3;
18     char broiled = 'c';
19     void *where_prohibited = NULL;
20
21     second_function(imidate);
22     imidate = 10;
23 }
24
25
26 void second_function(int a)
27 {
28     int b = a;
29 }
```

Frame for `main()`

Frame for `first_function()`

Return to `main()`, line 9

Storage space for an int

Storage space for a char

Storage space for a void *

Frame for `second_function()`:

Return to `first_function()`, line 22

Storage space for an int

Storage for the int parameter named `a`

Analyzing the Stack in GDB

(gdb) backtrace|bt

- shows the call trace (call stack)
- without function calls:
 - #0 `main()` at `program.c:10`
 - one frame on the stack, numbered 0, and it belongs to `main()`
- after call to function `display()`
 - #0 `display (z=5, zptr=0xbffffb34)` at `program.c:15` #1 `0x08048455` in `main()` at `program.c:10`

- two stack frames
 - frame 1 belonging to main()
 - frame 0 belonging to display()
- each frame listing gives
 - the arguments to that function
 - the line number that's currently being executed within that frame

```
1 (gdb) info frame
2 # displays info about the current stack frame, including its return address and saved
  register values
3
4 (gdb) info locals
5 # list the local variables of the function corresponding to the stack frame, with their
  current values
6
7 (gdb) info args
8 # list the argument values of the corresponding function call
```

Other Useful Commands

```
1 (gdb) info functions
2 # list all functions in the program
3
4 (gdb) list
5 # lists source code lines around the current line
```

Lab 5

- Download old version of coreutils with buggy ls program
 - Untar, configure, make
- Bug: ls -t mishandles files whose time stamps are very far in the past. It seems to act as if they are in the future

```
$ touch -d '1918-11-11 11:00 GMT' wwi-armistice
$ touch now
$ sleep 1
$ touch now1
$ ls -lt wwi-armistice now now1
```

Output:

```
-rw-r--r-- 1 eggert eggert 0 Nov 11 1918 wwi-armistice
-rw-r--r-- 1 eggert eggert 0 Feb 5 15:57 now1
-rw-r--r-- 1 eggert eggert 0 Feb 5 15:57 now
```

- **Reproduce the Bug**
 - Follow steps on lab web page
- **Simplify input**
 - Run ls with -l and -t options only
- **Debug**
 - Use gdb to figure out what's wrong
 - \$ gdb ./ls
 - (gdb) run -lt wwi-armistice now now1
(run from the directory where the compiled ls lives)
- **Patch**
 - Construct a patch "lab5.diff" containing your fix
 - It should contain a ChangeLog entry followed by the output of diff -u

- **Reproduce the Bug**
 - Follow steps on lab web page
- **Simplify input**
 - Run ls with -l and -t options only
- **Debug**
 - Use gdb to figure out what's wrong
 - \$ gdb ./ls
 - (gdb) run -lt wwi-armistice now now1
(run from the directory where the compiled ls lives)
- **Patch**
 - Construct a patch "lab5.diff" containing your fix
 - It should contain a ChangeLog entry followed by the output of diff -u
- Don't forget to answer all questions! (lab5.txt)
- Make sure not to submit a reverse patch! (lab5.diff)
- "Try to reproduce the problem in your home directory, instead of the \$tmp directory. How well does SEASnet do?"
 - Timestamps represented as seconds since Unix Epoch
 - SEASnet NFS filesystem has unsigned 32-bit time stamps
 - Local File System on Linux server has signed 32-bit time stamps
 - If you touch the files on the NFS filesystem it will return timestamp around 2054
 - => files have to be touched on local filesystem (df -l)
- Use "info functions" to look for relevant starting point
- Compiler optimizations: -O2 -> -O0
 - ./configure CFLAGS="...-O0"