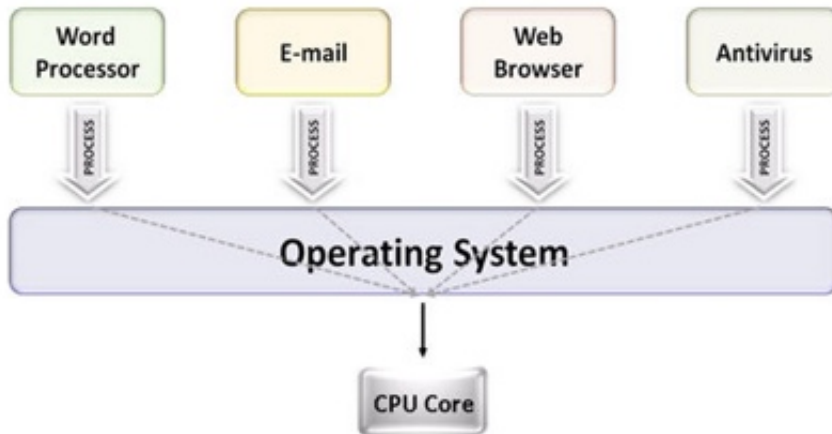


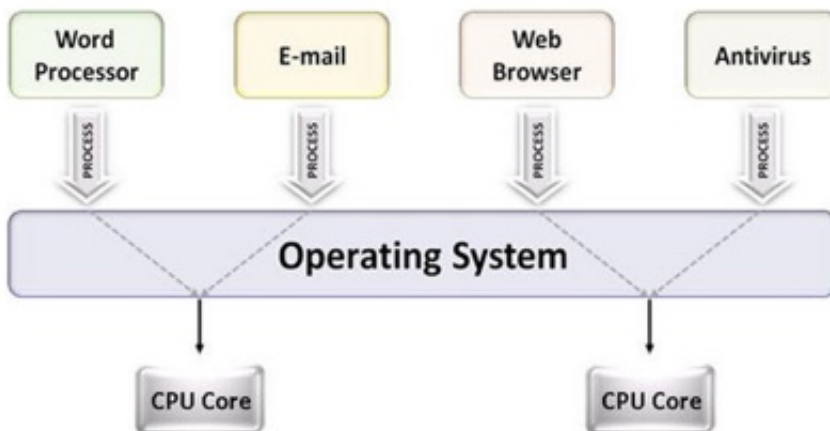
Multiprocessing

- use of multiple CPUs/cores to run multiple tasks simultaneously

Uniprocessing System



Multiprocessing System




Parallelism

- executing several computations simultaneously to gain performance
- different forms of parallelism
 - **multitasking**
 - several processes are scheduled alternately or possibly simultaneously on a multiprocessing system
 - **multithreading**
 - same job is broken logically into pieces (threads) which may be executed simultaneously on

What is a thread?

- a flow of instructions, path of execution within a process
- the smallest unit of processing scheduled by OS
- a process consists of *at least one* thread
- multiple threads can be run on:
 - **a uniprocessor (time-sharing)**
 - processor switches between different threads
 - switch back and forth, giving each thread a certain amount of time or CPU clock cycles to run
 - parallelism is an illusion

Multiple
threads
sharing a
single CPU



- **a multiprocessor**
 - multiple processors or cores run the threads at the same time
 - true parallelism

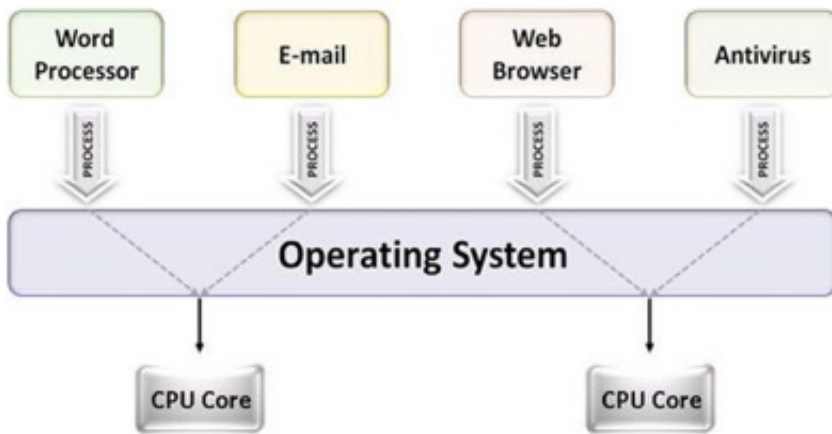
Multiple
threads on
multiple
CPUs



Multitasking vs. Multithreading

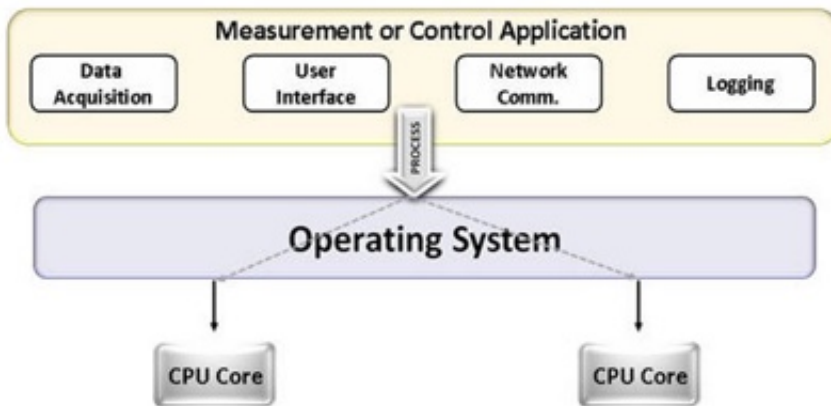
Multitasking

- multiple applications using multiple cores




Multithreading

- single application using multiple cores for different threads



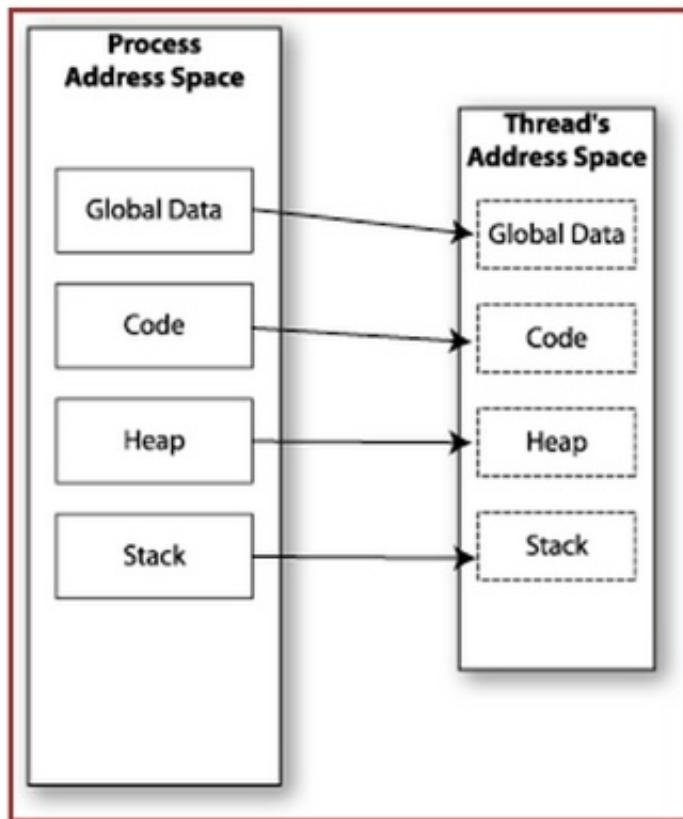
Multithreading

Multitasking

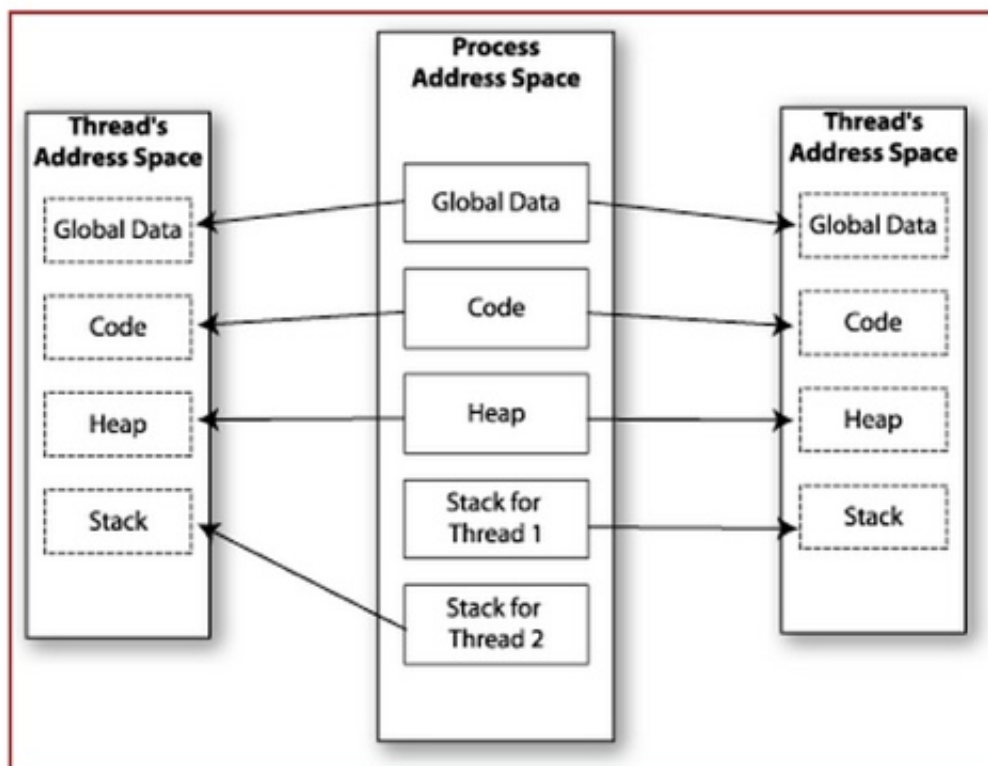


Memory Layout

Single-Threaded Program



Multithreaded Program



Multitasking

- `$ tr -cs 'A-Za-z' '[\n*]' sort -u | comm -23 - words`

- process 1 (tr)
- process 2 (sort)
- process 3 (comm)
- each process has its own address space
- how do these processes communicate?
 - pipes/system calls

Multithreading


- threads share all of the process's memory except for their stacks
- data sharing requires no extra work (no syscalls, pipes, etc.)

Shared Memory

- makes multithreaded programming...
 - **powerful**
 - can easily access data and share it among threads
 - **more efficient**
 - no need for syscalls when sharing data
 - thread creation and destruction less expensive than process creation and destruction
 - **non-trivial**
 - have to prevent several threads from accessing and changing the same shared data at the same time (synchronization)

Race Condition

```
1 int count = 0;
2 void increment() {
3     count = count + 1;
4 }
5
6 int main() {
7     increment();
8     increment();
9     increment();
10 }
```



- if we multithread this program across 3 threads, how the OS schedules the calls to increment makes a huge difference
 - if **thread C** executes between **thread B**'s read from count and its write to count, **thread B** will incorrectly "increment" the value of count to the same number it already was
- **synchronization needed**
 - each thread will "lock" each variable it modifies, preventing other threads from modifying the variable while it is being modified by the original thread
 - the thread will "unlock" that variable when it is done with it, freeing the variable to be modified by other threads

Multithreading & Multitasking: Comparison

Multithreading

- threads share same address space
 - light-weight creation/destruction
 - easy inter-thread communication
 - an error in one thread can bring down all threads in process
 - e.g. accessing a nullptr in one thread brings down entire process

Multitasking

- processes are insulated from each other
 - expensive creation/destruction
 - expensive IPC (inter-process communication)
 - an error in one process cannot bring down another process

- but results may not be as expected

Lab 8

- evaluate the performance of multithreaded sort
- add `/usr/local/cs/bin` to `PATH`
 - `$ export PATH=/usr/local/cs/bin:$PATH`
- generate a file containing 10M random **double-precision floating point numbers**, one per line with no white space
 - `/dev/urandom`: pseudo-random number generator
- `od`
 - write the contents of its input files to stdout in a user-specified format
 - options
 - `-t f`: double-precision floating point
 - `-N <count>`: format no more than count bytes of input
- `sed, tr`
 - remove address, delete spaces, add newlines between each float
- use `time -p` to time the command `sort -g` on the data you generated
- send output to `/dev/null`
- run `sort` with the `--parallel` option and the
 - `-g` option: compare by general numerical value
 - use `time` command to record the real, user and system time when running `sort` with 1, 2, 4, 8 threads

```
1 $ time -p sort -g file_name > /dev/null (1 thread)
2 $ time -p sort -g --parallel=[2, 4, or 8] file_name > /dev/null
```

- record the times and steps in `log.txt`